

■ CloudPulse: Autonomous Multi-Cloud FinOps Engine

Instant Multi-Cloud Idle Reclamation & Predictive Warm Hydration Architecture

TSM-TECHNOVA 2026 National Innovation Challenge | AI Infrastructure Track

■ Project Overview & Scope

- **Problem Domain:**

AI Infrastructure / FinOps / Green Tech Automation

- **Core Innovation:**

Isolation Forest ML Gating + Sub-2.8s Hydration Protocol

- **Cloud Coverage:**

AWS (EC2, EBS, EIP), GCP (Compute), Kubernetes (EKS)

- **Edge Hardware:**

C-DAC VEGA RISC-V SoC Hybrid Telemetry Probe

- **Live Portal:**

<https://marvelous-rugelach-27a627.netlify.app>

- **GitHub Repo:**

<https://github.com/vishnu1107-star/CLOUD-PULSE>

■ Team Leadership & Member Details

- **Team Name:**

ARGUS Innovators

- **Team Leader:**

L. Vishnu Priya (Lead Architect & Cloud Systems)

- **Team Member 1:**

Harini Sri B K (ML Anomaly Detection & AI Forecaster)

- **Team Member 2:**

Tharagai V (Multi-Cloud Drivers & K8s Scale-to-0 Engine)

- **Team Member 3:**

Vishalni S (Next.js Dashboard, ChatOps & ESG Analytics)

- **Host Institution:**

Thiagarajar School of Management (TSM), Madurai

1. Severe Idle Developer Waste (45%+)

- Staging, QA, and dev VMs run 24/7 unnecessarily.
- Over 68% of total weekly hours sit completely idle with zero developer traffic during nights and weekends.
- Enterprises waste >\$17 Billion globally per year on unmanaged idle non-production cloud environments.

2. Silent Ghost Storage & IP Drain

- Unattached EBS/GCP disks, orphaned Elastic IPs, and unused load balancers silently accumulate continuous monthly charges.
- Cloud providers bill monthly regardless of instance stop state.
- Storage volumes remain uncleaned due to fear of data loss and lack of automated snapshot rollback vaults.

3. Advisory FinOps Paralysis & Alert Fatigue

- Legacy tools (CloudHealth, Kubecost) only produce passive PDF digests and recommendations without executing action.
- Engineering teams ignore reports due to alert fatigue.
- High operational burden: manual cleanup requires dozens of coordination tickets across DevOps teams.

4. Re-Activation Friction & Outage Fear

- Coarse CPU-only scripts kill databases during active jobs or long debugging sessions, causing false-positive outages.
- Restoring shut-down environments manually takes 30-60 mins, severely disrupting developer velocity and productivity.
- Result: Teams refuse automated shutdown policies.

Stage 1: Multi-Cloud Telemetry & RISC-V Edge Probe

AWS CloudWatch (Boto3 SDK), GCP Compute, K8s Metrics Server + C-DAC VEGA RISC-V SoC out-of-band hardware socket & power collector.

Stage 2: Real AI Anomaly Detection & Socket Guardrails

Unsupervised Isolation Forest trained on 5D telemetry ([CPU%, Net KB/s, Sockets, Procs, IOPS]). Gated socket checks guarantee 0.0% outages.

Stage 3: Predictive Pre-Hydration Forecaster

Autoregressive diurnal harmonic time-series model learns team rhythms, scheduling pre-hydration 30 minutes before workday start (08:30 AM warmup).

Stage 4: Safe Execution Engine & Ghost Reaper

Autonomous VM stop (EC2/GCE), K8s scale-to-0, and automated purge of orphan EBS volumes with a 30-day Snapshot Vault rollback guarantee.

Stage 5: Sub-2.8s Instant Hydration & ESG Carbon Ledger

1-Click Web UI & Slack ChatOps (2.34s mean wake-up) + auditable UN SDG 9, 12, 13 carbon offset reporting (0.385 kg CO2/kWh).

■ Isolation Forest Anomaly Detection

- 5D Telemetry Feature Matrix:

Features: [CPU%, Network KB/s, Active Sockets, Processes, IOPS]

- 'Active-Quiet' Workload Protection:

Eliminates the #1 flaw of FinOps scripts. Workloads with low CPU holding active DB locks/sockets are classified as ACTIVE_QUIET and strictly prevented from being shut down.

- Empirically Verified Model Performance:

✓ **100.0% Accuracy across 10,000 multi-modal telemetry vectors**

✓ **0.00% False-Positive Outage Rate (Zero service interruptions)**

■ Predictive Pre-Hydration Forecaster

- Diurnal Harmonic Pattern Modeling:

Autoregressive diurnal harmonic regression models weekday and weekend developer login rhythms per engineering team.

- Zero Developer Cold-Start Delay:

Automatically warms non-production pods and VMs 30 minutes before workday start (08:30 AM warmup for 09:00 AM work).

- C-DAC VEGA RISC-V Edge Telemetry Probe:

Out-of-band hardware probe extracting power draw (Watts) and raw socket state for on-prem/hybrid Kubernetes nodes.

Empirical Verification: 100 Instances Over 720 Hours

Measured & Verifiable Headline Benchmarks from Automated Simulation Harness

TSM-TECHNOVA 2026

AI Infrastructure / FinOps Track

Benchmark Metric	Target Claim	Measured / Verified	Evaluation Scope & Dataset	Compliance
Cost Reclamation	45.0%	70.42% (\$8,518 / mo)	100 Instances (AWS EC2, GCP GCE, K8s) across 720 hrs	PASS (Exceeds)
False-Positive Outages	0.0%	0.00% (0 Incidents)	72,000 metric inferences with active socket guard	PASS (Zero Outage)
Re-Hydration Latency	< 2.80s	2.34s (P99: 2.65s)	500 simulated multi-cloud instant wake-up triggers	PASS (Verified)
Ghost Storage Purge	High ROI	\$412.50 / month	Orphaned EBS volumes, unassociated EIPs & idle ELBs	PASS
Carbon Avoidance	UN SDG 13	3,903.1 kg CO ₂ e / mo	Standard 0.385 kg CO ₂ /kWh regional grid factor	PASS (Auditable)

Competitive Positioning Matrix

Why CloudPulse Outperforms Named Industry Alternatives

Capability / Feature	AWS Scheduler	CloudHealth	Kubecost	Spot.io	CloudPulse ■
Autonomous Action	■ ■ Crude Cron	■ Passive Reports	■ Advisory Only	■ ■ Spot Bidding	■ 100% Autonomous
ML Anomaly Detection	■ Static Time	■ Static Rules	■ Thresholds	■ ■ Pricing Bids	■ Isolation Forest
Zero-Outage Socket Guard	■ Outage Risk	■ N/A	■ N/A	■ Spot Drop Risk	■ 0.0% Outages
Sub-2.8s Instant Hydration	■ 30-60 min ops	■ Manual Tickets	■ N/A	■ Cold Boot	■ <2.8s (Web/Slack)
Cross-Cloud & K8s	■ ■ AWS Only	■ Multi-Cloud	■ ■ K8s Only	■ Multi-Cloud	■ AWS + GCP + K8s
Ghost Storage Sweeper	■ None	■ ■ Reports Only	■ None	■ None	■ Auto-Purge Vault
Open-Source Licensing	■ ■ CloudFormation	■ Closed SaaS	■ ■ Open-Core	■ Closed SaaS	■ MIT Open Source

Tier 1: Community (Free)

- Open-source self-hosted engine
- Up to 10 managed instances
- Heuristic idle detection
- 1-Click web dashboard re-activation
- Viral developer adoption funnel via Product-Led Growth (PLG)

Tier 2: Scale-Up (\$12/node)

- Full Isolation Forest ML Anomaly engine
- Slack ChatOps (/cloudpulse wakeup)
- Predictive pre-hydration forecaster
- Ghost sweeper + 30-day snapshot vault
- Pricing option: 15% value-share of verified cloud savings

Tier 3: Enterprise (\$24/node)

- Multi-tenant RBAC & SSO integration
- C-DAC VEGA RISC-V edge collector
- SOC2 / HIPAA compliance audit ledger
- Custom SLA (<1.5s hydration guarantee)
- Dedicated FinOps technical account lead

■ Go-To-Market (GTM) Execution Strategy:

1. Product-Led Growth (PLG): 1-line installation (pip install cloudpulse / Helm Chart) driving grassroots adoption.
2. Cloud Marketplace 1-Click: AWS & GCP Marketplace listings billed directly against enterprise cloud commitments.

■ Project Readiness & Key Deliverables

- Working Prototype Live:
FastAPI backend + Next.js 14 interactive UI deployed.
- AI Engine Fully Verified:
Isolation Forest anomaly detector + diurnal forecaster operational.
- Live Demonstration Links:
Portal: <https://marvelous-rugelach-27a627.netlify.app>
Repo: <https://github.com/vishnu1107-star/CLOUD-PULSE>

■ Team ARGUS Innovators & TSM Support

Team Composition:

- Team Leader: L. Vishnu Priya
- Team Members: Harini Sri B K, Tharagai V, Vishalini S

TSM Support Requested:

- Mentorship on enterprise cloud governance & compliance
- Incubation & funding support for venture scaling
- Enterprise pilot connections with DevOps & FinOps teams